

# 1. Epoch

- **Definition:** An epoch refers to one complete pass through the entire training dataset during the training process.
- **Explanation:** In machine learning, especially when using **gradient-based optimization algorithms** (like SGD or Adam), the model's weights are updated after each batch of data. A **single epoch** means that the algorithm has seen and learned from every sample in the training dataset once. Typically, you train for many epochs to allow the model to learn better from the data.
- **Example:** If you have 1000 samples in your dataset and you train for 10 epochs, it means the model will pass over all 1000 samples 10 times.

# 2. Batch Size

- **Definition:** The batch size refers to the number of samples processed together in one forward and backward pass during the training of the model.
- **Explanation:** Rather than passing the entire dataset through the model at once (which can be computationally expensive), the data is split into smaller subsets called **batches**. For each batch, the model computes the predictions, calculates the loss, and updates the weights. The **batch size** determines how many samples are used in one such update.
- **Example:** If your dataset has 1000 samples and you set the batch size to 100, then the model will process 100 samples at a time and update its weights after each batch. In this case, it will require 10 steps to complete one epoch.

# 3. Samples Iteration (Steps per Epoch)

- **Definition:** The number of iterations or **steps per epoch** refers to how many times the model will update its weights during one epoch, which is determined by the **batch size** and the total number of samples.
- **Explanation:** If you have a dataset with  $N$  samples and a batch size of  $B$ , the number of iterations (or steps) per epoch is calculated as:

$$\text{Iterations per epoch} = \frac{N}{B}$$

- **Example:** If you have 1000 samples and your batch size is 100, you will have 10 iterations per epoch. After each iteration, the model will update the weights based on the gradient computed from that batch.

## 4. Weight

- **Definition:** Weights are the parameters in a machine learning model that are learned during training. These weights determine the strength of the connection between input features and output predictions.
- **Explanation:** When a neural network (or any other model) is trained, the goal is to adjust the weights so that the model's predictions become more accurate. During training, the model uses an optimization algorithm (like gradient descent) to adjust the weights based on the error or **loss**.
- **Example:** In a simple linear regression model  $y = w_1x_1 + w_2x_2 + b$ , the weights  $w_1$  and  $w_2$  are adjusted during training to minimize the error between the predicted and actual values.

## 5. Accuracy

- **Definition:** Accuracy is a metric used to evaluate how well the model is performing, usually by measuring the percentage of correct predictions.
- **Explanation:** Accuracy is the ratio of the number of correct predictions to the total number of predictions. In classification problems, it's the number of times the predicted label matches the true label divided by the total number of samples.
- **Example:** If a model correctly predicts the label for 90 out of 100 samples, the accuracy is 90%.

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}} \times 100$$

- **Note:** Accuracy may not always be the best metric, especially for imbalanced datasets where one class is more frequent than others. In such cases, other metrics like **precision**, **recall**, or **F1-score** might provide a better measure of performance.

## 6. Loss

- **Definition:** Loss is a function that measures how far off the model's predictions are from the actual values. It quantifies the error or discrepancy between predicted values and the actual target values.
- **Explanation:** The goal during training is to **minimize** the loss function. The loss provides feedback to the model on how well it's doing during training, and based on this feedback, the weights are updated to improve the model.

- **Example:** In **regression** problems, the **mean squared error (MSE)** is commonly used as a loss function:

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Where  $y_i$  is the true value and  $\hat{y}_i$  is the predicted value.

- In **classification** problems, a common loss function is **cross-entropy loss** (or **log loss**), which measures the difference between the predicted probabilities and the actual class labels.

## How They Work Together:

- During training, the model iterates over the data in **epochs**, processing the data in **batches**. In each batch, the model computes the predictions, compares them with the actual values, and calculates the **loss**.
- Based on the **loss**, the model updates the **weights** using an optimization algorithm (like **stochastic gradient descent**).
- After each epoch, the **accuracy** is computed to evaluate how well the model is performing on the training data.

## Example Workflow:

Let's put this into perspective with a simple neural network training process:

### 1. Epoch 1:

- Model processes the first batch of data (100 samples, given a batch size of 100) and computes the predictions.
- The model calculates the **loss** (e.g., MSE or cross-entropy) between predicted and true values.
- Based on the **loss**, the model updates its **weights**.
- The **accuracy** is evaluated for the epoch (how many correct predictions were made).

### 2. Epoch 2:

- The process is repeated with the next batch of data.
- After all epochs, the **final accuracy** and **loss** will indicate how well the model has learned from the data.

## In Summary:

- **Epoch:** One complete pass over the entire training dataset.
- **Batch Size:** Number of samples processed before the model updates its weights.
- **Iterations per Epoch:** Number of updates the model will make during an epoch (determined by batch size).
- **Weights:** Parameters that the model learns and updates during training.
- **Accuracy:** The percentage of correct predictions the model makes.
- **Loss:** The error or difference between the predicted and true values, which the model tries to minimize.